

ETL Project Report

Title: Toronto Neighbourhood Population Census vs Crime Data

Finding Data (Extract)

Data Sources:

- **Toronto Neighbourhood 2016 Census of Population**

https://ckan0.cf.opendata.inter.prod-toronto.ca/en_AU/dataset/neighbourhood-profiles

Data Format – CSV file

In this csv file all the data was formatted in an unstructured manner where all the neighbourhood was aligned in columns and all the various population data was aligned in rows as below.

| A | B | C | D | E | F | G | H | I | J |
|-----|---------------|---------------------|-----------------|---|-----------------|----------------|----------------|----------------|----------------|
| _id | Category | Topic | Data Source | Characteristic | City of Toronto | Agincourt | Agincourt | Alderwood | Annex |
| 1 | Neighbourhood | Neighbourhood | City of Toronto | Neighbourhood Number | | 129 | 128 | 20 | 95 |
| 2 | Neighbourhood | Neighbourhood | City of Toronto | TSNS2020 Designation | | No Designation | No Designation | No Designation | No Designation |
| 3 | Population | Population | Census Profile | Population, 2016 | 2,731,571 | 29,113 | 23,757 | 12,054 | 30,526 |
| 4 | Population | Population | Census Profile | Population, 2011 | 2,615,060 | 30,279 | 21,988 | 11,904 | 29,177 |
| 5 | Population | Population | Census Profile | Population Change 2011-2016 | 4.50% | -3.90% | 8.00% | 1.30% | 4.60% |
| 6 | Population | Population | Census Profile | Total private dwellings | 1,179,057 | 9,371 | 8,535 | 4,732 | 18,109 |
| 7 | Population | Population | Census Profile | Private dwellings occupied by usual residents | 1,112,929 | 9,120 | 8,136 | 4,616 | 15,934 |
| 8 | Population | Population | Census Profile | Population density per square kilometre | 4,334 | 3,929 | 3,034 | 2,435 | 10,863 |
| 9 | Population | Population | Census Profile | Land area in square kilometres | 630.2 | 7.41 | 7.83 | 4.95 | 2.81 |
| 10 | Population | Age characteristics | Census Profile | Children (0-14 years) | 398,135 | 3,840 | 3,075 | 1,760 | 2,360 |
| 11 | Population | Age characteristics | Census Profile | Youth (15-24 years) | 340,270 | 3,705 | 3,360 | 1,235 | 3,750 |
| 12 | Population | Age characteristics | Census Profile | Working Age (25-54 years) | 1,229,555 | 11,305 | 9,965 | 5,220 | 15,040 |
| 13 | Population | Age characteristics | Census Profile | Pre-retirement (55-64 years) | 336,670 | 4,230 | 3,265 | 1,825 | 3,480 |
| 14 | Population | Age characteristics | Census Profile | Seniors (65+ years) | 426,945 | 6,045 | 4,105 | 2,015 | 5,910 |
| 15 | Population | Age characteristics | Census Profile | Older Seniors (85+ years) | 66,000 | 925 | 555 | 320 | 1,040 |

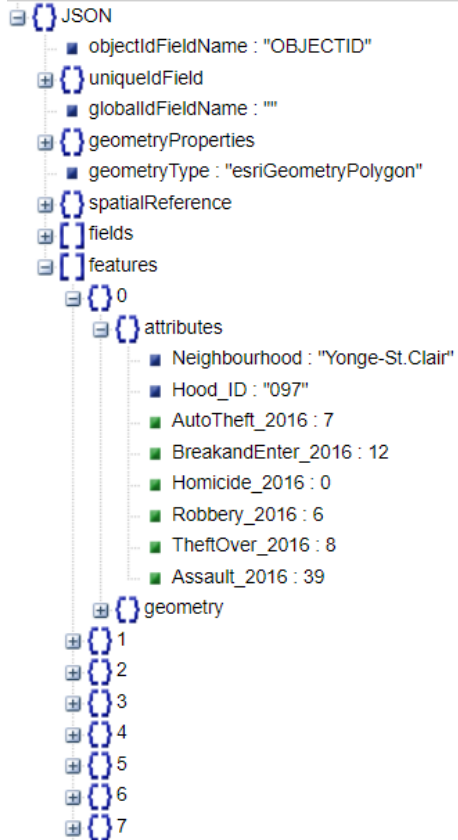
- **Toronto Neighbourhood 2016 Crime Data**

http://data.torontopolice.on.ca/datasets/neighbourhood-crime-rates-boundary-file-/geoservice?geometry=-80.421%2C43.542%2C-78.335%2C43.890&orderBy=Hood_ID

Data Format – GeoJSON file

From this link, we opted for 2016 crime data for each Toronto neighbourhood in GeoJSON format. In JSON viewer, the data structure looked like below.

Viewer Text



Data Cleanup and & Analysis (Transform)

From the 2016 Toronto population data, all the neighbourhood names and its hood id and age population spectrum data were extracted. In original data, the neighbourhood name was the column index, hood id was in row format and the population spectrum was various row format structure shown as below.

Jupyter census_crime_df Last Checkpoint: Last Saturday at 7:30 PM (autosaved) Python 3

```
In [1]: import pandas as pd
        from sqlalchemy import create_engine

In [2]: toronto_profile_csvpath = "Resources/neighbourhood-profiles-2016-csv.csv"

        toronto_2016_raw_df = pd.read_csv(toronto_profile_csvpath)
        toronto_2016_raw_df.head(5)
```

Out[2]:

| | _id | Category | Topic | Data Source | Characteristic | City of Toronto | Agincourt North | Agincourt South-Malvern West | Alderwood | Annex | ... | Willowdale West | Willowridge-Martingrove-Richview |
|---|-----|---------------------------|---------------------------|--------------------------------|-----------------------------|-----------------|-----------------|------------------------------|----------------|----------------|-----|-----------------|----------------------------------|
| 0 | 1 | Neighbourhood Information | Neighbourhood Information | City of Toronto | Neighbourhood Number | NaN | 129 | 128 | 20 | 95 | ... | 37 | 7 |
| 1 | 2 | Neighbourhood Information | Neighbourhood Information | City of Toronto | TSNS2020 Designation | NaN | No Designation | No Designation | No Designation | No Designation | ... | No Designation | No Designation |
| 2 | 3 | Population | Population and dwellings | Census Profile 98-316-X2016001 | Population, 2016 | 2,731,571 | 29,113 | 23,757 | 12,054 | 30,526 | ... | 16,936 | 22,156 |
| 3 | 4 | Population | Population and dwellings | Census Profile 98-316-X2016001 | Population, 2011 | 2,615,060 | 30,279 | 21,988 | 11,904 | 29,177 | ... | 15,004 | 21,343 |
| 4 | 5 | Population | Population and dwellings | Census Profile 98-316-X2016001 | Population Change 2011-2016 | 4.50% | -3.90% | 8.00% | 1.30% | 4.60% | ... | 12.90% | 3.80% |

5 rows x 146 columns

The unwanted columns were dropped, and the specific columns and rows were extracted and put in the individual dataframes by using iloc function and the dataframes were merged to create a single neighbourhood population dataframe as below

```
In [3]: df = toronto_2016_raw_df
        df = df.drop(columns=['_id', 'Category', 'Data Source', 'Characteristic', 'Topic', 'City of Toronto', 'Topic'])
        dfT = df.T
        dfTi = dfT.iloc[:,0:3]
        dfT2 = dfT.iloc[:,9:15]
        dft_merge = dfTi.merge(dfT2,how='outer',left_index=True,right_index=True)
        dft_merge = dft_merge.drop(columns=[1])
        dft_merge
```

Out[3]:

| | 0 | 2 | 9 | 10 | 11 | 12 | 13 | 14 |
|------------------------------|-----|--------|-------|-------|--------|-------|-------|-------|
| Agincourt North | 129 | 29,113 | 3,840 | 3,705 | 11,305 | 4,230 | 6,045 | 925 |
| Agincourt South-Malvern West | 128 | 23,757 | 3,075 | 3,360 | 9,965 | 3,265 | 4,105 | 555 |
| Alderwood | 20 | 12,054 | 1,760 | 1,235 | 5,220 | 1,825 | 2,015 | 320 |
| Annex | 95 | 30,526 | 2,360 | 3,750 | 15,040 | 3,480 | 5,910 | 1,040 |
| Banbury-Don Mills | 42 | 27,695 | 3,605 | 2,730 | 10,810 | 3,555 | 6,975 | 1,640 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| Wychwood | 94 | 14,349 | 1,860 | 1,320 | 6,420 | 1,595 | 3,150 | 880 |
| Yonge-Eglinton | 100 | 11,817 | 1,800 | 1,225 | 5,860 | 1,325 | 1,600 | 165 |
| Yonge-St.Clair | 97 | 12,528 | 1,210 | 920 | 5,960 | 1,540 | 2,905 | 470 |
| York University Heights | 27 | 27,593 | 4,045 | 4,750 | 12,290 | 2,965 | 3,530 | 400 |
| Yorkdale-Glen Park | 31 | 14,804 | 1,960 | 1,870 | 5,860 | 1,810 | 3,295 | 775 |

140 rows x 8 columns

As columns transformed into rows and rows into columns, we had to name all the individual columns indexes its appropriate names.

```

dft_merge.columns = ['Neighbourhood ID', 'Population', 'Children (0-14 years)', 'Youth (15-24 years)', 'Working Age (25-54 years)', 'Pre-retirement (55-64 years)', 'Seniors (65+ years)', 'Older Seniors (85+ years)']
dft_merge = dft_merge.reset_index()
dft_merge = dft_merge.rename(columns={'Neighbourhood ID': 'id', "index": "Neighbourhood", "Pre-retirement (55-64 years)": "Pre_retirement (55-64 years)"})
dft_fin = dft_merge.set_index("id")
dft_fin

```

[4]:

| | Neighbourhood | Population | Children (0-14 years) | Youth (15-24 years) | Working Age (25-54 years) | Pre_retirement (55-64 years) | Seniors (65+ years) | Older Seniors (85+ years) |
|-----|------------------------------|------------|-----------------------|---------------------|---------------------------|------------------------------|---------------------|---------------------------|
| id | | | | | | | | |
| 129 | Agincourt North | 29,113 | 3,840 | 3,705 | 11,305 | 4,230 | 6,045 | 925 |
| 128 | Agincourt South-Malvern West | 23,757 | 3,075 | 3,360 | 9,965 | 3,265 | 4,105 | 555 |
| 20 | Alderwood | 12,054 | 1,760 | 1,235 | 5,220 | 1,825 | 2,015 | 320 |
| 95 | Annex | 30,526 | 2,360 | 3,750 | 15,040 | 3,480 | 5,910 | 1,040 |
| 42 | Banbury-Don Mills | 27,695 | 3,605 | 2,730 | 10,810 | 3,555 | 6,975 | 1,640 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 94 | Wychwood | 14,349 | 1,860 | 1,320 | 6,420 | 1,595 | 3,150 | 880 |
| 100 | Yonge-Eglinton | 11,817 | 1,800 | 1,225 | 5,860 | 1,325 | 1,600 | 165 |
| 97 | Yonge-St.Clair | 12,528 | 1,210 | 920 | 5,960 | 1,540 | 2,905 | 470 |
| 27 | York University Heights | 27,593 | 4,045 | 4,750 | 12,290 | 2,965 | 3,530 | 400 |
| 31 | Yorkdale-Glen Park | 14,804 | 1,960 | 1,870 | 5,860 | 1,810 | 3,295 | 775 |

140 rows × 8 columns

Database transformation from “jupyter notebook” to “pgAdmin” in sql format requires the keys and values of the source databases to be presented in the format as pgAdmin would if the database tables are created with the program’s query tools, failure to format properly would result in pgAdmin unable to load the source data.

A few steps were taken to clean the column keys (column titles);

1. no symbols, such as hyphen, plus sign and parentheses
2. no space, which we had to replace with underscore
3. no capital letters, all capital letters in column title has to changed into lower case (code used: `df.columns=map(str.lower,df.columns)`)

In addition, In order for pgAdmin to process dataframe numbers as integers, during the cleaning process, we had to make sure, I.E 3,711 is expressed as 3711, removing any delimiter.

```

### data cleaning, removing all the symls progres can't process and rename titles to lowercase

cols = dft_fin.columns
cols = cols.map(lambda x: x.replace(' ', '_') if isinstance(x, (str, str)) else x)
cols = cols.map(lambda x: x.replace('(', '_') if isinstance(x, (str, str)) else x)
cols = cols.map(lambda x: x.replace(')', '_') if isinstance(x, (str, str)) else x)
cols = cols.map(lambda x: x.replace('-', '_to_') if isinstance(x, (str, str)) else x)
cols = cols.map(lambda x: x.replace('+', 'plus') if isinstance(x, (str, str)) else x)
dft_fin.columns = cols

dft_fin.columns = map(str.lower, dft_fin.columns)
cleaned_census_df = dft_fin

cleaned_census_df

```

[4]:

| | neighbourhood | population | children_0_to_14_years | youth_15_to_24_years | working_age_25_to_54_years | pre_retirement_55_to_64_years | seniors_65pl |
|-----|------------------------------|------------|------------------------|----------------------|----------------------------|-------------------------------|--------------|
| id | | | | | | | |
| 129 | Agincourt North | 29113 | 3840 | 3705 | 11305 | 4230 | |
| 128 | Agincourt South-Malvern West | 23757 | 3075 | 3360 | 9965 | 3265 | |

```

dft_fin['Population'] = dft_fin['Population'].str.replace(',', '').astype(int)
dft_fin['Children_0_to_14_years'] = dft_fin['Children_0_to_14_years'].str.replace(',', '').astype(int)
dft_fin['Youth_15_to_24_years'] = dft_fin['Youth_15_to_24_years'].str.replace(',', '').astype(int)
dft_fin['Working_Age_25_to_54_years'] = dft_fin['Working_Age_25_to_54_years'].str.replace(',', '').astype(int)
dft_fin['Pre_retirement_55_to_64_years'] = dft_fin['Pre_retirement_55_to_64_years'].str.replace(',', '').astype(int)
dft_fin['Seniors_65plus_years'] = dft_fin['Seniors_65plus_years'].str.replace(',', '').astype(int)
dft_fin['Older_Seniors_85plus_years'] = dft_fin['Older_Seniors_85plus_years'].str.replace(',', '').astype(int)
dft_fin

```

5]:

| | Neighbourhood | Population | Children_0_to_14_years | Youth_15_to_24_years | Working_Age_25_to_54_years | Pre_retirement_55_to_64_years | Seniors_65pl |
|-----|------------------------------|------------|------------------------|----------------------|----------------------------|-------------------------------|--------------|
| id | | | | | | | |
| 129 | Agincourt North | 29113 | 3840 | 3705 | 11305 | 4230 | |
| 128 | Agincourt South-Malvern West | 23757 | 3075 | 3360 | 9965 | 3265 | |
| 20 | Alderwood | 12054 | 1760 | 1235 | 5220 | 1825 | |
| 95 | Annex | 30526 | 2360 | 3750 | 15040 | 3480 | |

For the downloaded data for Toronto Neighbourhood 2016 crime was in GeoJSON format. It seemed very unstructured compared to usual JSON format, shown below:

```

{
  "type": "FeatureCollection",
  "features": [
    {
      "type": "Feature",
      "properties": {
        "Neighbourhood": "Yonge-St.Clair",
        "Hood_ID": "097",
        "AutoTheft_2016": 7,
        "BreakandEnter_2016": 12,
        "Homicide_2016": 0,
        "Robbery_2016": 6,
        "TheftOver_2016": 8,
        "Assault_2016": 39
      },
      "geometry": {
        "type": "Polygon",
        "coordinates": [
          [
            [
              [
                -79.38,
                43.65,
                -79.37,
                43.65,
                -79.37,
                43.66,
                -79.38,
                43.66,
                -79.38,
                43.65
              ]
            ]
          ]
        ]
      }
    }
  ]
}

```

From this data, the attributes dictionary was the only data needed. Everything else was noise.

Here ‘features’ list contains ‘attributes’ data and ‘Neighbourhood’, ‘Hood_ID’, etc. are properties of ‘attributes’ dictionary.

In order to extract the targeted data, json was imported, json.load function was used and all the required data was appended within the individual lists.

```

import requests
import json

with open('Resources/Neighbourhood_Crime_Rates_Boundary_File_.geojson') as f:
    data = json.load(f)

Neighbourhood = []
Hood_ID = []
Assault_2016 = []
AutoTheft_2016 = []
BreakandEnter_2016 = []
Homicide_2016 = []
Robbery_2016 = []
TheftOver_2016 = []

for feature in data['features']:
    #print(feature['geometry']['type'])
    #print(feature['geometry']['coordinates'])
    #print(feature)
    #print(feature['properties']['Neighbourhood'])
    #print(feature['properties']['Hood_ID'])

    Neighbourhood.append(feature['properties']['Neighbourhood'])
    Hood_ID.append(feature['properties']['Hood_ID'])
    Assault_2016.append(feature['properties']['Assault_2016'])
    AutoTheft_2016.append(feature['properties']['AutoTheft_2016'])
    BreakandEnter_2016.append(feature['properties']['BreakandEnter_2016'])
    Homicide_2016.append(feature['properties']['Homicide_2016'])
    Robbery_2016.append(feature['properties']['Robbery_2016'])
    TheftOver_2016.append(feature['properties']['TheftOver_2016'])

```

Toronto_crime dataframe was created with all the data extracted within the individual lists.

```

toronto_crime_df = pd.DataFrame(
    {
        "Neighbourhood": Neighbourhood,
        "id": Hood_ID,
        "Assault_2016": Assault_2016,
        "AutoTheft_2016": AutoTheft_2016,
        "BreakandEnter_2016": BreakandEnter_2016,
        "Homicide_2016": Homicide_2016,
        "Robbery_2016": Robbery_2016,
        "TheftOver_2016": TheftOver_2016
    }
)
toronto_crime_df.head()

```

| | Neighbourhood | id | Assault_2016 | AutoTheft_2016 | BreakandEnter_2016 | Homicide_2016 | Robbery_2016 | TheftOver_2016 |
|---|-------------------------|-----|--------------|----------------|--------------------|---------------|--------------|----------------|
| 0 | Yonge-St.Clair | 097 | 39 | 7 | 12 | 0 | 6 | 8 |
| 1 | York University Heights | 027 | 361 | 105 | 98 | 2 | 70 | 37 |
| 2 | Lansing-Westgate | 038 | 68 | 27 | 41 | 0 | 9 | 5 |
| 3 | Yorkdale-Glen Park | 031 | 174 | 41 | 66 | 1 | 24 | 26 |
| 4 | Stonegate-Queensway | 016 | 76 | 12 | 49 | 0 | 16 | 4 |

As mentioned before, column index names were converted to lowercase as PGAdmin was being difficult accepting this table as is.

```

crime_df.columns = map(str.lower, crime_df.columns)
cleaned_crime_df = crime_df

cleaned_crime_df

```

```

In [ ]:
neighbourhood assault_2016 autotheft_2016 breakandenter_2016 homicide_2016 robbery_2016 theftover_2016
id
097 Yonge-St.Clair 39 7 12 0 6 8
027 York University Heights 361 105 98 2 70 37
038 Lansing-Westgate 68 27 41 0 9 5
031 Yorkdale-Glen Park 174 41 66 1 24 26
016 Stonegate-Queensway 76 12 49 0 16 4
...
130 Milliken 88 35 67 0 40 10
046 Pleasant View 58 8 32 0 13 3
094 Wychwood 86 18 29 0 9 4
056 Leaside-Bennington 33 4 32 0 8 2
108 Briar Hill-Belgravia 64 16 31 0 21 4

```

Connection to the SQL postgres as below.

```

In [ ]:
connection_string = "*****:*****@localhost:5432/toronto_tb"
engine = create_engine(f'postgresql://{connection_string}')

```

```

In [ ]:
engine.table_names()

```

```

Out[6]: ['crime', 'census']

```

```

In [ ]:
cleaned_census_df.to_sql(name='census', con=engine, if_exists='append', index=True)

```

```

In [ ]:
cleaned_crime_df.to_sql(name='crime', con=engine, if_exists='append', index=True)

```

Loading to Database (Load)

Among the two different data sources, there is a common key which the data sources each used to identify census area, which is the Toronto neighbourhood number (neighbourhood ID, referred simply as 'ID' in the final table).

Data base was created in pgAdmin and tables were created as below:

```
1  CREATE TABLE census
2  ( id INT PRIMARY KEY,
3      Neighbourhood VARCHAR,
4      Population INT ,
5      Children_0_to_14_years INT,
6      Youth_15_to_24_years INT,
7      Working_Age_25_to_54_years INT,
8      Pre_retirement_55_to_64_years INT,
9      Seniors_65plus_years INT,
10     Older_Seniors_85plus_years INT
11 );
12
13 CREATE TABLE crime (
14 id INT PRIMARY KEY,
15     Neighbourhood VARCHAR,
16     Assault_2016 INT,
17     AutoTheft_2016 INT,
18     BreakandEnter_2016 INT,
19     Homicide_2016 INT,
20     robbery_2016 INT,
21     TheftOver_2016 INT
22 );
23
24 drop table census
25 drop table crime
26
```

The databases are loaded through jupyter's notebook and then connected and transferred to pgAdmin for easier merging and browsing.

The Neighbourhood Profile database contains data of population in each neighbourhood and their respective age demographic, whilst the Toronto Crime Rates database contains data of the count of different crimes committed in different neighbourhood.

SQL queries:

```

1  SELECT * FROM crime;
2
3  SELECT * FROM census;
4
5  -- Join tables on id
6  SELECT census.id, crime.neighbourhood, census.population
7  FROM census
8  INNER JOIN crime
9  ON census.id = crime.id
10 ORDER BY id ASC;
11

```

Toronto Neighbourhood Profiles:

| | id [PK] integer | neighbourhood character varying | assault_2016 integer | autotheft_2016 integer | breakandenter_2016 integer | homicide_2016 integer | robbery_2016 integer | theftover_2016 integer |
|----|--------------------|------------------------------------|-------------------------|---------------------------|-------------------------------|--------------------------|-------------------------|---------------------------|
| 1 | 1 | West Humber-Clairville | 307 | 321 | 131 | 4 | 100 | 41 |
| 2 | 2 | Mount Olive-Silverston... | 270 | 43 | 34 | 1 | 102 | 4 |
| 3 | 3 | Thistletown-Beaumont... | 39 | 13 | 23 | 0 | 16 | 1 |
| 4 | 4 | Rexdale-Kipling | 77 | 22 | 16 | 0 | 17 | 0 |
| 5 | 5 | Elms-Old Rexdale | 61 | 16 | 10 | 0 | 18 | 0 |
| 6 | 6 | Kingsview Village-The ... | 125 | 42 | 34 | 2 | 24 | 5 |
| 7 | 7 | Willowridge-Martingrov... | 102 | 49 | 35 | 0 | 57 | 4 |
| 8 | 8 | Humber Heights-West... | 30 | 16 | 23 | 1 | 8 | 4 |
| 9 | 9 | Edenbridge-Humber Va... | 30 | 29 | 46 | 0 | 7 | 2 |
| 10 | 10 | Princess-Rosethorn | 30 | 24 | 33 | 1 | 32 | 3 |

Toronto Neighbourhood Crime Rates:

| | id [PK] integer | neighbourhood character varying | population integer | children_0_to_14_years integer | youth_15_to_24_years integer | working_age_25_to_54_years integer |
|----|--------------------|------------------------------------|-----------------------|-----------------------------------|---------------------------------|---------------------------------------|
| 1 | 1 | West Humber-Clairville | 33312 | 5060 | 5445 | 13845 |
| 2 | 2 | Mount Olive-Silverston... | 32954 | 7090 | 5240 | 13615 |
| 3 | 3 | Thistletown-Beaumont... | 10360 | 1730 | 1410 | 4160 |
| 4 | 4 | Rexdale-Kipling | 10529 | 1640 | 1355 | 4300 |
| 5 | 5 | Elms-Old Rexdale | 9456 | 1805 | 1440 | 3700 |
| 6 | 6 | Kingsview Village-The ... | 22000 | 4240 | 3020 | 8635 |
| 7 | 7 | Willowridge-Martingrov... | 22156 | 3555 | 2625 | 8140 |
| 8 | 8 | Humber Heights-West... | 10948 | 1450 | 1140 | 3790 |
| 9 | 9 | Edenbridge-Humber Va... | 15535 | 2120 | 1805 | 5940 |
| 10 | 10 | Princess-Rosethorn | 11051 | 1770 | 1580 | 3825 |

By merging the two tables together, we are hoping to discover; 1), which neighbourhood is more prominent in crime and 2) if there is a relation between crime rate and the age demographic.